

Project Title: Exploring Movie Success: Revenue, Ratings, and Genre Analysis

Business Understanding

Your company now sees all the big companies creating original video content and they want to get in on the fun. They have decided to create a new movie studio, but they don't know anything about creating movies. You are charged with exploring what types of films are currently doing the best at the box office. You must then translate those findings into actionable insights that the head of your company's new movie studio can use to help decide what type of films to create.

Introduction

The success of a movie is influenced by multiple factors, including genre, release timing, audience reception, and market performance. Understanding these factors can help studios make informed decisions to maximize both critical acclaim and financial returns. This analysis explores the relationships between movie ratings, genres, release months, and box office performance, aiming to provide actionable insights for strategic planning in movie production and distribution. By examining patterns in domestic and worldwide profits, as well as audience ratings, the study identifies key drivers of commercial success and offers recommendations to optimize future releases.

Problem Statement

Studios often struggle to predict which factors, such as genre, release timing, or ratings drive a movie's financial success. This study aims to identify the key drivers of profitability to guide better production and release decisions.

Objectives

1. To analyze how movie genre impacts domestic and worldwide box office profits.
2. To examine the effect of release month on movie profitability.
3. To investigate the relationship between audience ratings and box office success.
4. To determine whether domestic performance can predict worldwide earnings.
5. To provide actionable recommendations for optimizing movie production and release strategies.

Data Understanding

The dataset used in this analysis contains information about movies, including their genre, release date, audience ratings, domestic gross, and worldwide gross. It provides insights into both critical reception and financial performance, allowing for an examination of patterns and relationships that influence movie success. It contained various dataset file formats, tsvs, csvs and an sqlite database, which all contained various columns eg genre, release date, audience ratings, domestic gross, and worldwide gross just to mention a few. These multiple datasets enable analysis of correlations between ratings, genre, release timing, and profits, helping to identify the factors most critical to commercial success.

Importing all Libraries

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
```

```
import matplotlib.pyplot as plt
from scipy.stats import norm
from scipy.stats import binom
```

Exploratory Data Analysis (EDA)

Exploratory Data Analysis was conducted to understand patterns, trends, and relationships in the movie dataset.

In [2]:

```
# Reading into the bom csv files to clean and aggregate the columns we need for analysis
bom = pd.read_csv("bom.movie_gross.csv")

#Get a feel of how the data looks like
bom.columns
```

Out[2]:

```
Index(['title', 'studio', 'domestic_gross', 'foreign_gross', 'year'], dtype='object')
```

In [3]:

```
bom
```

Out[3]:

| | title | studio | domestic_gross | foreign_gross | year |
|------|---|------------|----------------|---------------|------|
| 0 | Toy Story 3 | BV | 415000000.0 | 652000000 | 2010 |
| 1 | Alice in Wonderland (2010) | BV | 334200000.0 | 691300000 | 2010 |
| 2 | Harry Potter and the Deathly Hallows Part 1 | WB | 296000000.0 | 664300000 | 2010 |
| 3 | Inception | WB | 292600000.0 | 535700000 | 2010 |
| 4 | Shrek Forever After | P/DW | 238700000.0 | 513900000 | 2010 |
| ... | ... | ... | ... | ... | ... |
| 3382 | The Quake | Magn. | 6200.0 | NaN | 2018 |
| 3383 | Edward II (2018 re-release) | FM | 4800.0 | NaN | 2018 |
| 3384 | EI Pacto | Sony | 2500.0 | NaN | 2018 |
| 3385 | The Swan | Synergetic | 2400.0 | NaN | 2018 |
| 3386 | An Actor Prepares | Grav. | 1700.0 | NaN | 2018 |

3387 rows × 5 columns

In [4]:

```
bom.describe()
```

Out[4]:

| | domestic_gross | year |
|-------|----------------|-------------|
| count | 3.359000e+03 | 3387.000000 |
| mean | 2.874585e+07 | 2013.958075 |
| std | 6.698250e+07 | 2.478141 |
| min | 1.000000e+02 | 2010.000000 |
| 25% | 1.200000e+05 | 2012.000000 |
| 50% | 1.400000e+06 | 2014.000000 |
| 75% | 2.790000e+07 | 2016.000000 |
| max | 9.367000e+08 | 2018.000000 |

```
In [5]:
```

```
bom.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   title            3387 non-null    object  
 1   studio           3382 non-null    object  
 2   domestic_gross   3359 non-null    float64 
 3   foreign_gross    2037 non-null    object  
 4   year             3387 non-null    int64  
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB
```

```
In [6]:
```

```
bom['foreign_gross'] = pd.to_numeric(bom['foreign_gross'], errors='coerce')
```

```
In [7]:
```

```
bom['foreign_gross'].astype(float)
```

```
Out[7]:
```

```
0      652000000.0
1      691300000.0
2      664300000.0
3      535700000.0
4      513900000.0
...
3382      NaN
3383      NaN
3384      NaN
3385      NaN
3386      NaN
Name: foreign_gross, Length: 3387, dtype: float64
```

```
In [8]:
```

```
bom.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   title            3387 non-null    object  
 1   studio           3382 non-null    object  
 2   domestic_gross   3359 non-null    float64 
 3   foreign_gross    2032 non-null    float64 
 4   year             3387 non-null    int64  
dtypes: float64(2), int64(1), object(2)
memory usage: 132.4+ KB
```

```
In [9]:
```

```
#Check missing rows in the bomdf
bom.isna().sum()
```

```
Out[9]:
```

```
title          0
studio         5
domestic_gross 28
foreign_gross  1355
year           0
dtype: int64
```

```
In [10]:
```

```
bom['foreign_gross'] = bom['foreign_gross'].fillna(bom['foreign_gross'].mean())
bom['domestic_gross'] = bom['domestic_gross'].fillna(bom['domestic_gross'].mean())
```

In [11]:

```
bom.isna().sum()
```

Out[11]:

```
title      0
studio     5
domestic_gross  0
foreign_gross 0
year       0
dtype: int64
```

In [12]:

```
bom=bom.dropna()
```

In [13]:

```
bom['studio'].isna().sum()
```

Out[13]:

```
0
```

In [14]:

```
bom.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3382 entries, 0 to 3386
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   title            3382 non-null   object 
 1   studio           3382 non-null   object 
 2   domestic_gross   3382 non-null   float64
 3   foreign_gross    3382 non-null   float64
 4   year             3382 non-null   int64  
dtypes: float64(2), int64(1), object(2)
memory usage: 158.5+ KB
```

In [15]:

```
bom.duplicated()
```

Out[15]:

```
0      False
1      False
2      False
3      False
4      False
...
3382  False
3383  False
3384  False
3385  False
3386  False
Length: 3382, dtype: bool
```

In [16]:

```
bomdf = ['title', 'domestic_gross', 'foreign_gross']
bom = bom[bomdf]
```

In [17]:

```
bom
```

Out [17]:

| | title | domestic_gross | foreign_gross |
|------|---|----------------|---------------|
| 0 | Toy Story 3 | 415000000.0 | 6.520000e+08 |
| 1 | Alice in Wonderland (2010) | 334200000.0 | 6.913000e+08 |
| 2 | Harry Potter and the Deathly Hallows Part 1 | 296000000.0 | 6.643000e+08 |
| 3 | Inception | 292600000.0 | 5.357000e+08 |
| 4 | Shrek Forever After | 238700000.0 | 5.139000e+08 |
| ... | ... | ... | ... |
| 3382 | The Quake | 6200.0 | 7.505704e+07 |
| 3383 | Edward II (2018 re-release) | 4800.0 | 7.505704e+07 |
| 3384 | El Pacto | 2500.0 | 7.505704e+07 |
| 3385 | The Swan | 2400.0 | 7.505704e+07 |
| 3386 | An Actor Prepares | 1700.0 | 7.505704e+07 |

3382 rows × 3 columns

In [18]:

```
rt_movie = pd.read_table("rt.movie_info.tsv")
```

In [19]:

```
rt_movie.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               1560 non-null    int64  
 1   synopsis         1498 non-null    object  
 2   rating           1557 non-null    object  
 3   genre            1552 non-null    object  
 4   director         1361 non-null    object  
 5   writer           1111 non-null    object  
 6   theater_date    1201 non-null    object  
 7   dvd_date         1201 non-null    object  
 8   currency          340 non-null    object  
 9   box_office        340 non-null    object  
 10  runtime          1530 non-null    object  
 11  studio           494 non-null    object  
dtypes: int64(1), object(11)
memory usage: 146.4+ KB
```

In [20]:

```
rt_movie.isna().sum()
```

Out [20]:

```
id                  0
synopsis          62
rating              3
genre              8
director          199
writer             449
theater_date      359
dvd_date           359
currency          1220
box_office         1220
runtime             30
studio             1066
dtype: int64
```

```
In [21]:
```

```
rt_movie['runtime'] = rt_movie['runtime'].str.replace('minutes', '', regex=False)
rt_movie['runtime']
```

```
Out[21]:
```

```
0      104
1      108
2      116
3      128
4      200
...
1555    106
1556     88
1557    111
1558    101
1559     94
Name: runtime, Length: 1560, dtype: object
```

```
In [22]:
```

```
rt_movie['runtime'] = pd.to_numeric(rt_movie['runtime'], errors='coerce')
```

```
In [23]:
```

```
rt_movie.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   id               1560 non-null   int64  
 1   synopsis         1498 non-null   object  
 2   rating            1557 non-null   object  
 3   genre             1552 non-null   object  
 4   director          1361 non-null   object  
 5   writer            1111 non-null   object  
 6   theater_date     1201 non-null   object  
 7   dvd_date          1201 non-null   object  
 8   currency          340 non-null   object  
 9   box_office        340 non-null   object  
 10  runtime           1530 non-null   float64 
 11  studio            494 non-null   object  
dtypes: float64(1), int64(1), object(10)
memory usage: 146.4+ KB
```

```
In [24]:
```

```
rt_movie['runtime'].value_counts()
```

```
Out[24]:
```

```
runtime
90.0      72
95.0      66
100.0     51
93.0      47
96.0      43
..
154.0      1
166.0      1
33.0       1
54.0       1
290.0     1
Name: count, Length: 142, dtype: int64
```

```
In [25]:
```

```
rt_movie['runtime'] = rt_movie['runtime'].fillna(rt_movie['runtime'].mode()[0])
```

```
rt_movie['runtime'].value_counts()
```

Out [25]:

```
runtime
90.0      102
95.0       66
100.0      51
93.0       47
96.0       43
...
154.0       1
166.0       1
33.0        1
54.0        1
290.0      1
Name: count, Length: 142, dtype: int64
```

In [26]:

```
rt_movie.isna().sum()
```

Out [26]:

```
id                  0
synopsis           62
rating              3
genre               8
director            199
writer              449
theater_date        359
dvd_date            359
currency            1220
box_office          1220
runtime              0
studio              1066
dtype: int64
```

In [27]:

```
rt_movie.head()
```

Out [27]:

| | | | | id | synopsis | rating | genre | director | writer | theater_date | dvd_date | currency | box_o |
|---|---|--|----|--|---------------------|---------------------------------------|--------------|-----------------|--------|--------------|----------|----------|-------|
| 0 | 1 | This gritty, fast-paced, and innovative police... | R | Action and Adventure Classics Drama | William Friedkin | Ernest Tidyman | Oct 9, 1971 | Sep 25, 2001 | NaN | | | | |
| 1 | 3 | New York City, not- too-distant- future: Eric Pa... | R | Drama Science Fiction and Fantasy | David Cronenberg | David Cronenberg Don DeLillo | Aug 17, 2012 | Jan 1, 2013 | \$ | 600 | | | |
| 2 | 5 | Illeana Douglas delivers a superb performance ... | R | Drama Musical and Performing Arts | Allison Anders | Allison Anders | Sep 13, 1996 | Apr 18, 2000 | NaN | | | | |
| 3 | 6 | Michael Douglas runs afoul of a treacherous su... | R | Drama Mystery and Suspense | Barry Levinson | Paul Attanasio Michael Crichton | Dec 9, 1994 | Aug 27, 1997 | NaN | | | | |
| 4 | 7 | NaN | NR | Drama Romance | Rodney Bennett | Giles Cooper | NaN | NaN | NaN | | | | |

In [28]:

```
rt_reviews = pd.read_table("rt.reviews.tsv", encoding="latin1")
rt_reviews.head()
```

Out[28]:

| id | review | rating | fresh | critic | top_critic | publisher | date |
|-----------|--|---------------|--------------|----------------|-------------------|------------------|-------------------|
| 0 3 | A distinctly gallows take on contemporary financial... n... | 3/5 | fresh | PJ Naborro | 0 | Patrick Naborro | November 10, 2018 |
| 1 3 | It's an allegory in search of a meaning that n... | NaN | rotten | Annalee Newitz | 0 | io9.com | May 23, 2018 |
| 2 3 | ... life lived in a bubble in financial dealin... | NaN | fresh | Sean Axmaker | 0 | Stream on Demand | January 4, 2018 |
| 3 3 | Continuing along a line introduced in last yea... | NaN | fresh | Daniel Kasman | 0 | MUBI | November 16, 2017 |
| 4 3 | ... a perverse twist on neorealism... | NaN | fresh | NaN | 0 | Cinema Scope | October 12, 2017 |

In [29]:

```
rt_reviews.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54432 entries, 0 to 54431
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          54432 non-null   int64  
 1   review       48869 non-null   object  
 2   rating       40915 non-null   object  
 3   fresh        54432 non-null   object  
 4   critic       51710 non-null   object  
 5   top_critic   54432 non-null   int64  
 6   publisher    54123 non-null   object  
 7   date         54432 non-null   object  
dtypes: int64(2), object(6)
memory usage: 3.3+ MB
```

In [30]:

```
rt_reviews["rating"] = pd.to_numeric(rt_reviews["rating"], errors='coerce').astype("float64")
```

In [31]:

```
rt_reviews.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 54432 entries, 0 to 54431
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   id          54432 non-null   int64  
 1   review       48869 non-null   object  
 2   rating       750 non-null    float64 
 3   fresh        54432 non-null   object  
 4   critic       51710 non-null   object  
 5   top_critic   54432 non-null   int64  
 6   publisher    54123 non-null   object  
 7   date         54432 non-null   object  
dtypes: float64(1), int64(2), object(5)
memory usage: 3.3+ MB
```

In [32]:

```
rt_reviews.isna().sum()
```

```
In [32]:
```

```
id          0
review      5563
rating     53682
fresh        0
critic     2722
top_critic    0
publisher    309
date         0
dtype: int64
```

```
In [33]:
```

```
rt_reviews['top_critic'].value_counts()
```

```
Out[33]:
```

```
top_critic
0    41336
1    13096
Name: count, dtype: int64
```

```
In [34]:
```

```
rt_reviews_1 = pd.read_table("rt.reviews.tsv", encoding="latin1")
```

```
In [35]:
```

```
rt_reviews_1 = rt_reviews_1.dropna()
```

```
In [36]:
```

```
rt_reviews_1.isna().sum()
```

```
Out[36]:
```

```
id          0
review      0
rating      0
fresh        0
critic      0
top_critic   0
publisher    0
date         0
dtype: int64
```

```
In [37]:
```

```
rt_reviews_1.duplicated()
```

```
Out[37]:
```

```
0      False
6      False
7      False
11     False
12     False
...
54419    False
54420    False
54421    False
54422    False
54424    False
Length: 33988, dtype: bool
```

```
In [38]:
```

```
rt_reviews_1.head()
```

```
Out[38]:
```

| id | review | rating | fresh | critic | top_critic | publisher | date |
|----|--------|--------|-------|--------|------------|-----------|------|
|----|--------|--------|-------|--------|------------|-----------|------|

| id | review | rating | fresh | critic | top_critic | publisher | date |
|-----------|---|---------------|--------------|----------------|-------------------|--------------------|-------------------|
| 0 3 | A distinctly gallows take on contemporary fina... | 3/5 | fresh | PJ Nabarro | 0 | Patrick Nabarro | November 10, 2018 |
| 6 3 | Quickly grows repetitive and tiresome, meander... | C | rotten | Eric D. Snider | 0 | EricDSnider.com | July 17, 2013 |
| 7 3 | Cronenberg is not a director to be daunted by ... | 2/5 | rotten | Matt Kelemen | 0 | Las Vegas CityLife | April 21, 2013 |
| 11 3 | While not one of Cronenberg's stronger films, ... | B- | fresh | Emanuel Levy | 0 | EmanuelLevy.Com | February 3, 2013 |
| 12 3 | Robert Pattinson works mighty hard to make Cos... | 2/4 | rotten | Christian Toto | 0 | Big Hollywood | January 15, 2013 |

In [39]:

```
df_tmdb = pd.read_csv("tmdb.movies.csv")
df_tmdb.columns
```

Out [39]:

```
Index(['Unnamed: 0', 'genre_ids', 'id', 'original_language', 'original_title',
       'popularity', 'release_date', 'title', 'vote_average', 'vote_count'],
      dtype='object')
```

In [40]:

```
df_tmdb
```

Out [40]:

| Unnamed: 0 | genre_ids | id | original_language | original_title | popularity | release_date | title | vote_ave |
|---------------|---------------------|--------|-------------------|--|------------|--------------|--|----------|
| 0 0 | [12, 14, 10751] | 12444 | en | Harry Potter and the Deathly Hallows: Part 1 | 33.533 | 2010-11-19 | Harry Potter and the Deathly Hallows: Part 1 | |
| 1 1 | [14, 12, 16, 10751] | 10191 | en | How to Train Your Dragon | 28.734 | 2010-03-26 | How to Train Your Dragon | |
| 2 2 | [12, 28, 878] | 10138 | en | Iron Man 2 | 28.515 | 2010-05-07 | Iron Man 2 | |
| 3 3 | [16, 35, 10751] | 862 | en | Toy Story | 28.005 | 1995-11-22 | Toy Story | |
| 4 4 | [28, 878, 12] | 27205 | en | Inception | 27.920 | 2010-07-16 | Inception | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 26512 26512 | [27, 18] | 488143 | en | Laboratory Conditions | 0.600 | 2018-10-13 | Laboratory Conditions | |
| 26513 26513 | [18, 53] | 485975 | en | _EXHIBIT_84xxx_ | 0.600 | 2018-05-01 | _EXHIBIT_84xxx_ | |
| 26514 26514 | [14, 28, 12] | 381231 | en | The Last One | 0.600 | 2018-10-01 | The Last One | |
| 26515 26515 | [10751, 12, 28] | 366854 | en | Trailer Made | 0.600 | 2018-06-22 | Trailer Made | |
| 26516 26516 | [53, 27] | 309885 | en | The Church | 0.600 | 2018-10-05 | The Church | |

26517 rows × 10 columns

In [41]:

```
df_tmdb.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26517 entries, 0 to 26516
Data columns (total 10 columns):
 #   Column          Non-Null Count  Dtype  
 0   id              26517 non-null   int64  
 1   genre_ids       26517 non-null   list    
 2   original_language 26517 non-null   object 
 3   original_title  26517 non-null   object 
 4   popularity      26517 non-null   float64
 5   release_date    26517 non-null   datetime64[ns]
 6   title           26517 non-null   object 
 7   vote_average    26517 non-null   float64
 8   vote_count      26517 non-null   int64  
 9   
```

```
# Column           Non-Null Count   Dtype  
---  -- 
0   Unnamed: 0      26517 non-null    int64  
1   genre_ids       26517 non-null    object  
2   id              26517 non-null    int64  
3   original_language 26517 non-null    object  
4   original_title   26517 non-null    object  
5   popularity       26517 non-null    float64 
6   release_date     26517 non-null    object  
7   title            26517 non-null    object  
8   vote_average     26517 non-null    float64 
9   vote_count       26517 non-null    int64  
dtypes: float64(2), int64(3), object(5)
memory usage: 2.0+ MB
```

In [42]:

```
# check for duplicates
df_tmdb.duplicated().sum()
```

Out[42]:

```
0
```

In [43]:

```
df_tmdb['release_date'] = pd.to_datetime(df_tmdb['release_date']).dt.month
```

In [44]:

```
df_tmdb["genre_ids"].value_counts().sort_index(ascending=False)
```

Out[44]:

```
genre_ids
[]                      2479
[99]                     3700
[99, 99]                  2
[99, 99, 99]               1
[99, 9648]                 4
...
[10402, 10751, 14, 10770, 35] 1
[10402, 10749]               3
[10402, 10749, 35]             2
[10402, 10749, 35, 18]          3
[10402, 10749, 18]              2
Name: count, Length: 2477, dtype: int64
```

In [45]:

```
df_tmdb = df_tmdb.rename(columns={'release_date':'release_month'})
```

In [46]:

```
tmdbd = ['title', 'vote_average', 'vote_count', 'release_month']
df_tmdb = df_tmdb[tmdbd]
```

In [47]:

```
df_tmdb
```

Out[47]:

| | | title | vote_average | vote_count | release_month |
|----------|---|-------|--------------|--------------|---------------|
| 0 | Harry Potter and the Deathly Hallows: Part 1 | | 7.7 | 10788 | 11 |
| 1 | How to Train Your Dragon | | 7.7 | 7610 | 3 |
| 2 | Iron Man 2 | | 6.8 | 12368 | 5 |
| 3 | Toy Story | | 7.9 | 10174 | 11 |
| 4 | Inception | | 8.3 | 22186 | 7 |

| | | title | vote_average | vote_count | release_month |
|-------|--|-----------------------|--------------|------------|---------------|
| 26512 | | Laboratory Conditions | 0.0 | 1 | 10 |
| 26513 | | _EXHIBIT_84xxx_ | 0.0 | 1 | 5 |
| 26514 | | The Last One | 0.0 | 1 | 10 |
| 26515 | | Trailer Made | 0.0 | 1 | 6 |
| 26516 | | The Church | 0.0 | 1 | 10 |

26517 rows × 4 columns

In [48]:

```
df_tn = pd.read_csv("tn.movie_budgets.csv")
df_tn
```

Out[48]:

| | | | id | release_date | movie | production_budget | domestic_gross | worldwide_gross |
|------|-----|--------------|----|--------------|---|-------------------|----------------|-----------------|
| 0 | 1 | Dec 18, 2009 | | | Avatar | \$425,000,000 | \$760,507,625 | \$2,776,345,279 |
| 1 | 2 | May 20, 2011 | | | Pirates of the Caribbean: On Stranger Tides | \$410,600,000 | \$241,063,875 | \$1,045,663,875 |
| 2 | 3 | Jun 7, 2019 | | | Dark Phoenix | \$350,000,000 | \$42,762,350 | \$149,762,350 |
| 3 | 4 | May 1, 2015 | | | Avengers: Age of Ultron | \$330,600,000 | \$459,005,868 | \$1,403,013,963 |
| 4 | 5 | Dec 15, 2017 | | | Star Wars Ep. VIII: The Last Jedi | \$317,000,000 | \$620,181,382 | \$1,316,721,747 |
| ... | ... | ... | | | ... | ... | ... | ... |
| 5777 | 78 | Dec 31, 2018 | | | Red 11 | \$7,000 | \$0 | \$0 |
| 5778 | 79 | Apr 2, 1999 | | | Following | \$6,000 | \$48,482 | \$240,495 |
| 5779 | 80 | Jul 13, 2005 | | | Return to the Land of Wonders | \$5,000 | \$1,338 | \$1,338 |
| 5780 | 81 | Sep 29, 2015 | | | A Plague So Pleasant | \$1,400 | \$0 | \$0 |
| 5781 | 82 | Aug 5, 2005 | | | My Date With Drew | \$1,100 | \$181,041 | \$181,041 |

5782 rows × 6 columns

In [49]:

```
df_tn.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5782 entries, 0 to 5781
```

```
Data columns (total 6 columns):
```

| # | Column | Non-Null Count | Dtype |
|---|-------------------|----------------|-----------------|
| 0 | id | 5782 | non-null int64 |
| 1 | release_date | 5782 | non-null object |
| 2 | movie | 5782 | non-null object |
| 3 | production_budget | 5782 | non-null object |
| 4 | domestic_gross | 5782 | non-null object |
| 5 | worldwide_gross | 5782 | non-null object |

```
dtypes: int64(1), object(5)
```

```
memory usage: 271.2+ KB
```

In [50]:

```
df_tn['domestic_gross'] = df_tn['domestic_gross'].str.replace('$', '', regex=False)
df_tn['worldwide_gross'] = df_tn['worldwide_gross'].str.replace('$', '', regex=False)
df_tn['production_budget'] = df_tn['production_budget'].str.replace('$', '', regex=False)
```

In [51]:

```
df_tn.head()
```

Out[51]:

In [51]:

| id | release_date | movie | production_budget | domestic_gross | worldwide_gross |
|-----------|-----------------------|---|--------------------------|-----------------------|------------------------|
| 0 | 1 Dec 18, 2009 | Avatar | 425,000,000 | 760,507,625 | 2,776,345,279 |
| 1 | 2 May 20, 2011 | Pirates of the Caribbean: On Stranger Tides | 410,600,000 | 241,063,875 | 1,045,663,875 |
| 2 | 3 Jun 7, 2019 | Dark Phoenix | 350,000,000 | 42,762,350 | 149,762,350 |
| 3 | 4 May 1, 2015 | Avengers: Age of Ultron | 330,600,000 | 459,005,868 | 1,403,013,963 |
| 4 | 5 Dec 15, 2017 | Star Wars Ep. VIII: The Last Jedi | 317,000,000 | 620,181,382 | 1,316,721,747 |

In [52]:

```
df_tn['production_budget'] = df_tn['production_budget'].str.replace('NA', '0', regex=False)
df_tn['domestic_gross'] = df_tn['domestic_gross'].str.replace('NA', '0', regex=False)
df_tn['worldwide_gross'] = df_tn['worldwide_gross'].str.replace('NA', '0', regex=False)
df_tn['production_budget'] = df_tn['production_budget'].str.replace(',', '', regex=False)
df_tn['domestic_gross'] = df_tn['domestic_gross'].str.replace(',', '', regex=False)
df_tn['worldwide_gross'] = df_tn['worldwide_gross'].str.replace(',', '', regex=False)
```

In [53]:

```
df_tn.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               5782 non-null    int64  
 1   release_date     5782 non-null    object 
 2   movie             5782 non-null    object 
 3   production_budget 5782 non-null   object 
 4   domestic_gross    5782 non-null    object 
 5   worldwide_gross   5782 non-null    object 
dtypes: int64(1), object(5)
memory usage: 271.2+ KB
```

In [54]:

```
df_tn['domestic_gross'] = pd.to_numeric(df_tn['domestic_gross'], errors='coerce').astype("float64")
df_tn['worldwide_gross'] = pd.to_numeric(df_tn['worldwide_gross'], errors='coerce').astype("float64")
df_tn['production_budget'] = pd.to_numeric(df_tn['production_budget'], errors='coerce').astype("float64")
```

In [55]:

```
df_tn.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5782 entries, 0 to 5781
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               5782 non-null    int64  
 1   release_date     5782 non-null    object 
 2   movie             5782 non-null    object 
 3   production_budget 5782 non-null   float64 
 4   domestic_gross    5782 non-null    float64 
 5   worldwide_gross   5782 non-null    float64 
dtypes: float64(3), int64(1), object(2)
memory usage: 271.2+ KB
```

In [56]:

```
df_tn = df_tn.rename(columns={'movie': 'title'})
```

In [57]:

```
tndf = ['title', 'production_budget', 'domestic_gross', 'worldwide_gross']
df_tn = df_tn[tndf]
```

In [58]:

```
df_tn
```

Out[58]:

| | | title | production_budget | domestic_gross | worldwide_gross |
|------|---|--------|-------------------|----------------|-----------------|
| 0 | | Avatar | 425000000.0 | 760507625.0 | 2.776345e+09 |
| 1 | Pirates of the Caribbean: On Stranger Tides | | 410600000.0 | 241063875.0 | 1.045664e+09 |
| 2 | Dark Phoenix | | 350000000.0 | 42762350.0 | 1.497624e+08 |
| 3 | Avengers: Age of Ultron | | 330600000.0 | 459005868.0 | 1.403014e+09 |
| 4 | Star Wars Ep. VIII: The Last Jedi | | 317000000.0 | 620181382.0 | 1.316722e+09 |
| ... | ... | ... | ... | ... | ... |
| 5777 | Red 11 | | 7000.0 | 0.0 | 0.000000e+00 |
| 5778 | Following | | 6000.0 | 48482.0 | 2.404950e+05 |
| 5779 | Return to the Land of Wonders | | 5000.0 | 1338.0 | 1.338000e+03 |
| 5780 | A Plague So Pleasant | | 1400.0 | 0.0 | 0.000000e+00 |
| 5781 | My Date With Drew | | 1100.0 | 181041.0 | 1.810410e+05 |

5782 rows × 4 columns

In [59]:

```
import sqlite3

conn = sqlite3.connect("im.db")

df2_sql = """
    SELECT name
    FROM SQLITE_MASTER
"""

pd.read_sql_query(df2_sql, conn)
```

Out[59]:

| | name |
|---|---------------|
| 0 | movie_basics |
| 1 | directors |
| 2 | known_for |
| 3 | movie_akas |
| 4 | movie_ratings |
| 5 | persons |
| 6 | principals |
| 7 | writers |

In [60]:

```
movie_ratingssql = """
    SELECT movie_basics.movie_id,
    movie_basics.genres, movie_basics.primary_title, movie_ratings.averageRating,
    movie_ratings.numVotes, movie_akas.region, movie_akas.language
"""
```

```

        FROM movie_basics
        JOIN movie_ratings ON movie_ratings.movie_id = movie_basics.movie_id
        JOIN movie_akas ON movie_akas.movie_id = movie_ratings.movie_id
        """
movie_ratingssql = pd.read_sql_query(movie_ratingssql, conn)

```

In [61]:

```
movie_ratingssql.to_csv('databasecsv.csv', index=False, encoding='utf-8')
```

In [62]:

```
dfdatabase = pd.read_csv('databasecsv.csv')
dfdatabase
```

Out[62]:

| | movie_id | genres | primary_title | averagerating | numvotes | region | language |
|--------|-----------|--------------------|---------------------------|---------------|----------|--------|----------|
| 0 | tt0063540 | Action,Crime,Drama | Sunghursh | 7.0 | 77 | NaN | NaN |
| 1 | tt0063540 | Action,Crime,Drama | Sunghursh | 7.0 | 77 | IN | NaN |
| 2 | tt0063540 | Action,Crime,Drama | Sunghursh | 7.0 | 77 | IN | hi |
| 3 | tt0063540 | Action,Crime,Drama | Sunghursh | 7.0 | 77 | IN | hi |
| 4 | tt0063540 | Action,Crime,Drama | Sunghursh | 7.0 | 77 | IN | hi |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 261801 | tt9905462 | Drama | Pengalila | 8.4 | 600 | IN | NaN |
| 261802 | tt9905462 | Drama | Pengalila | 8.4 | 600 | IN | en |
| 261803 | tt9911774 | Drama | Padmavyuhathile Abhimanyu | 8.4 | 365 | NaN | NaN |
| 261804 | tt9911774 | Drama | Padmavyuhathile Abhimanyu | 8.4 | 365 | IN | ml |
| 261805 | tt9911774 | Drama | Padmavyuhathile Abhimanyu | 8.4 | 365 | IN | ml |

261806 rows × 7 columns

In [63]:

```
dfdatabase = dfdatabase.drop_duplicates()
```

In [64]:

```
dfdatabase.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 242723 entries, 0 to 261804
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   movie_id    242723 non-null   object 
 1   genres       241653 non-null   object 
 2   primary_title 242723 non-null   object 
 3   averagerating 242723 non-null   float64
 4   numvotes     242723 non-null   int64  
 5   region        205843 non-null   object 
 6   language      35124 non-null   object 
dtypes: float64(1), int64(1), object(5)
memory usage: 14.8+ MB
```

In [65]:

```
dfdatabase.isna().sum()
```

Out[65]:

| | |
|---------------|------|
| movie_id | 0 |
| genres | 1070 |
| primary_title | 0 |

```
primary_title      0
averagerating     0
numvotes          0
region            36880
language          207599
dtype: int64
```

In [66]:

```
dfdatabasedrop = dfdatabase.dropna()
```

In [67]:

```
dfdatabasedrop.isna().sum()
```

Out[67]:

```
movie_id          0
genres            0
primary_title    0
averagerating    0
numvotes          0
region            0
language          0
dtype: int64
```

In [68]:

```
dfdatabasedrop.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 34992 entries, 2 to 261804
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   movie_id         34992 non-null   object 
 1   genres            34992 non-null   object 
 2   primary_title    34992 non-null   object 
 3   averagerating    34992 non-null   float64
 4   numvotes          34992 non-null   int64  
 5   region            34992 non-null   object 
 6   language          34992 non-null   object 
dtypes: float64(1), int64(1), object(5)
memory usage: 2.1+ MB
```

In [69]:

```
dfdatabasedrop
```

Out[69]:

| | movie_id | genres | primary_title | averagerating | numvotes | region | language |
|--------|-----------|----------------------|---------------------------------|---------------|----------|--------|----------|
| 2 | tt0063540 | Action,Crime,Drama | Sunghursh | 7.0 | 77 | IN | hi |
| 8 | tt0066787 | Biography,Drama | One Day Before the Rainy Season | 7.2 | 43 | XWW | en |
| 23 | tt0069204 | Comedy,Drama | Sabse Bada Sukh | 6.1 | 13 | IN | bn |
| 24 | tt0069204 | Comedy,Drama | Sabse Bada Sukh | 6.1 | 13 | IN | en |
| 29 | tt0100275 | Comedy,Drama,Fantasy | The Wandering Soap Opera | 6.5 | 119 | XWW | en |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 261791 | tt9899850 | Drama,Thriller | The Agitation | 4.9 | 14 | XWW | en |
| 261794 | tt9899860 | Drama,Thriller | Watching This Movie Is a Crime | 8.1 | 7 | XWW | en |
| 261796 | tt9899880 | Comedy | Columbus | 5.8 | 5 | IR | fa |
| 261802 | tt9905462 | Drama | Pengalila | 8.4 | 600 | IN | en |
| 261804 | tt9911774 | Drama | Padmavyuhathile Abhimanyu | 8.4 | 365 | IN | ml |

34992 rows × 7 columns

```
In [70]:
```

```
dfdatabasedrop = dfdatabasedrop.rename(columns={'primary_title':'title'})
```

```
In [71]:
```

```
bom.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3382 entries, 0 to 3386
Data columns (total 3 columns):
 #   Column      Non-Null Count  Dtype  
---  --  
 0   title       3382 non-null   object  
 1   domestic_gross  3382 non-null  float64 
 2   foreign_gross 3382 non-null  float64 
dtypes: float64(2), object(1)
memory usage: 105.7+ KB
```

```
In [72]:
```

```
themovie_df = (
    bom
        .merge(df_tmdb, on='title')
        .merge(df_tn, on='title')
        .merge(dfdatabasedrop, on='title')
)
#themovie_df = pd.merge(bom, df_tmdb, df_tn, dfdatabasedrop, on = 'title')
```

```
In [73]:
```

```
themovie_df
```

```
Out[73]:
```

| | | title | domestic_gross_x | foreign_gross | vote_average | vote_count | release_month | production_budget | domestic_gross |
|------|--------------|--------------|------------------|---------------|--------------|------------|---------------|-------------------|----------------|
| 0 | Toy Story 3 | Toy Story 3 | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| 1 | Toy Story 3 | Toy Story 3 | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| 2 | Toy Story 3 | Toy Story 3 | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| 3 | Toy Story 3 | Toy Story 3 | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| 4 | Toy Story 3 | Toy Story 3 | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6005 | Gotti | Gotti | 4300000.0 | 7.505704e+07 | 5.2 | 231 | 6 | 10000000.0 | 4286367. |
| 6006 | Mandy | Mandy | 1200000.0 | 7.505704e+07 | 3.5 | 2 | 1 | 6000000.0 | 1214525. |
| 6007 | Mandy | Mandy | 1200000.0 | 7.505704e+07 | 6.2 | 618 | 9 | 6000000.0 | 1214525. |
| 6008 | Lean on Pete | Lean on Pete | 1200000.0 | 7.505704e+07 | 6.9 | 133 | 4 | 8000000.0 | 1163056. |
| 6009 | Lean on Pete | Lean on Pete | 1200000.0 | 7.505704e+07 | 6.9 | 133 | 4 | 8000000.0 | 1163056. |

```
6010 rows × 15 columns
```

In [74]:

```
themovie_df = themovie_df.drop_duplicates()
```

In [75]:

```
themovie_df.value_counts().sum()
```

Out[75]:

5461

In [76]:

```
themovie_df
```

Out[76]:

| | | title | domestic_gross_x | foreign_gross | vote_average | vote_count | release_month | production_budget | domestic_gross_ |
|------|--------------|-------|------------------|---------------|--------------|------------|---------------|-------------------|-----------------|
| 0 | Toy Story 3 | | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| 1 | Toy Story 3 | | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| 2 | Toy Story 3 | | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| 3 | Toy Story 3 | | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| 4 | Toy Story 3 | | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6004 | Gotti | | 4300000.0 | 7.505704e+07 | 5.2 | 231 | 6 | 10000000.0 | 4286367. |
| 6005 | Gotti | | 4300000.0 | 7.505704e+07 | 5.2 | 231 | 6 | 10000000.0 | 4286367. |
| 6006 | Mandy | | 1200000.0 | 7.505704e+07 | 3.5 | 2 | 1 | 6000000.0 | 1214525. |
| 6007 | Mandy | | 1200000.0 | 7.505704e+07 | 6.2 | 618 | 9 | 6000000.0 | 1214525. |
| 6008 | Lean on Pete | | 1200000.0 | 7.505704e+07 | 6.9 | 133 | 4 | 8000000.0 | 1163056. |

5461 rows × 15 columns

In [77]:

```
themovie_dropcol = themovie_df.drop(columns= ['language', 'region'])
```

In [78]:

```
themovie_df2 = themovie_dropcol.drop_duplicates()
```

In [79]:

```
themovie_df2
```

Out[79]:

| | | title | domestic_gross_x | foreign_gross | vote_average | vote_count | release_month | production_budget | domestic_gro |
|---|-------------|-------|------------------|---------------|--------------|------------|---------------|-------------------|--------------|
| 0 | Toy Story 3 | | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 415004880. |

| 8 | Inception | domestic_gross_x | foreign_gross | vote_average | vote_count | release_month | production_budget | domestic_gro |
|------|-------------------------------------|------------------|---------------|--------------|------------|---------------|-------------------|--------------|
| 17 | Shrek Forever After | 238700000.0 | 5.139000e+08 | 6.1 | 3843 | 5 | 165000000.0 | 2387361 |
| 25 | The Twilight Saga: Eclipse | 300500000.0 | 3.980000e+08 | 6.0 | 4909 | 6 | 68000000.0 | 3005317 |
| 30 | Iron Man 2 | 312400000.0 | 3.115000e+08 | 6.8 | 12368 | 5 | 170000000.0 | 3124333 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6000 | Destroyer | 1500000.0 | 4.000000e+06 | 5.9 | 176 | 12 | 9000000.0 | 15333 |
| 6004 | Gotti | 4300000.0 | 7.505704e+07 | 5.2 | 231 | 6 | 10000000.0 | 42863 |
| 6006 | Mandy | 1200000.0 | 7.505704e+07 | 3.5 | 2 | 1 | 6000000.0 | 12145 |
| 6007 | Mandy | 1200000.0 | 7.505704e+07 | 6.2 | 618 | 9 | 6000000.0 | 12145 |
| 6008 | Lean on Pete | 1200000.0 | 7.505704e+07 | 6.9 | 133 | 4 | 8000000.0 | 11630 |

1335 rows × 13 columns

In [80]:

```
themovie_df2.to_csv('databaseforvisio.csv', index=False, encoding='utf-8')
```

In [81]:

```
themovie_df2['profit'] = themovie_df2['worldwide_gross'] - themovie_df2['production_budget']
themovie_df2['d_profit'] = themovie_df2['domestic_gross_y'] - themovie_df2['production_budget']
```

/tmp/ipykernel_21326/1521134713.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
themovie_df2['profit'] = themovie_df2['worldwide_gross'] - themovie_df2['production_budget']
```

/tmp/ipykernel_21326/1521134713.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
themovie_df2['d_profit'] = themovie_df2['domestic_gross_y'] - themovie_df2['production_budget']
```

In [82]:

```
themovie_df2
```

Out[82]:

| | title | domestic_gross_x | foreign_gross | vote_average | vote_count | release_month | production_budget | domestic_gro |
|----|---------------------------|------------------|---------------|--------------|------------|---------------|-------------------|--------------|
| 0 | Toy Story 3 | 415000000.0 | 6.520000e+08 | 7.7 | 8340 | 6 | 200000000.0 | 4150048 |
| 8 | Inception | 292600000.0 | 5.357000e+08 | 8.3 | 22186 | 7 | 160000000.0 | 2925761 |
| 17 | Shrek Forever After | 238700000.0 | 5.139000e+08 | 6.1 | 3843 | 5 | 165000000.0 | 2387361 |

The

| 25 | Twilight Saga: Eclipse | domestic_gross | foreign_gross | vote_average | vote_count | release_month | production_budget | domestic_box_office |
|------|------------------------|----------------|---------------|--------------|------------|---------------|-------------------|---------------------|
| 30 | Iron Man 2 | 312400000.0 | 3.115000e+08 | 6.8 | 12368 | 5 | 170000000.0 | 3124333 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6000 | Destroyer | 1500000.0 | 4.000000e+06 | 5.9 | 176 | 12 | 9000000.0 | 15333 |
| 6004 | Gotti | 4300000.0 | 7.505704e+07 | 5.2 | 231 | 6 | 10000000.0 | 42863 |
| 6006 | Mandy | 1200000.0 | 7.505704e+07 | 3.5 | 2 | 1 | 6000000.0 | 12145 |
| 6007 | Mandy | 1200000.0 | 7.505704e+07 | 6.2 | 618 | 9 | 6000000.0 | 12145 |
| 6008 | Lean on Pete | 1200000.0 | 7.505704e+07 | 6.9 | 133 | 4 | 8000000.0 | 11630 |

1335 rows × 15 columns

Statistical testing, insights extraction, and actionable recommendations.

What Genre of movie receives the highest ratings?

In [83]:

```
# Split and explode genres if not already done
themovie_df2['genres'] = themovie_df2['genres'].str.split(',')
themovie_df2 = themovie_df2.explode('genres')
themovie_df2['genres'] = themovie_df2['genres'].str.strip()
```

```
# Average rating by genre
genre_ratings = (
    themovie_df2.groupby('genres')['averagerating']
        .mean()
        .sort_values(ascending=False)
)
```

```
genre_ratings.head(10)
```

```
/tmp/ipykernel_21326/3119211118.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
themovie_df2['genres'] = themovie_df2['genres'].str.split(',')
```

Out[83]:

```
genres
Documentary      7.170588
History          7.155556
Biography         7.024031
Sport             6.947619
War               6.857143
Western           6.800000
Drama             6.690831
Animation         6.687912
Musical            6.633333
Sci-Fi             6.604386
Name: averagerating, dtype: float64
```

In [84]:

```
themovie_df2['genres'].value_counts(sort="descending")
```

Out[84]:

```
genres
Drama          698
Comedy         450
Action          389
Adventure       321
Thriller        253
Crime           212
Romance          88
Horror           50
Biography        29
Mystery          16
Sci-Fi            14
Fantasy          108
Animation         91
Family             71
History            45
Music              40
Sport                21
Documentary        17
War                  14
Western               8
Musical                 3
Name: count, dtype: int64
```

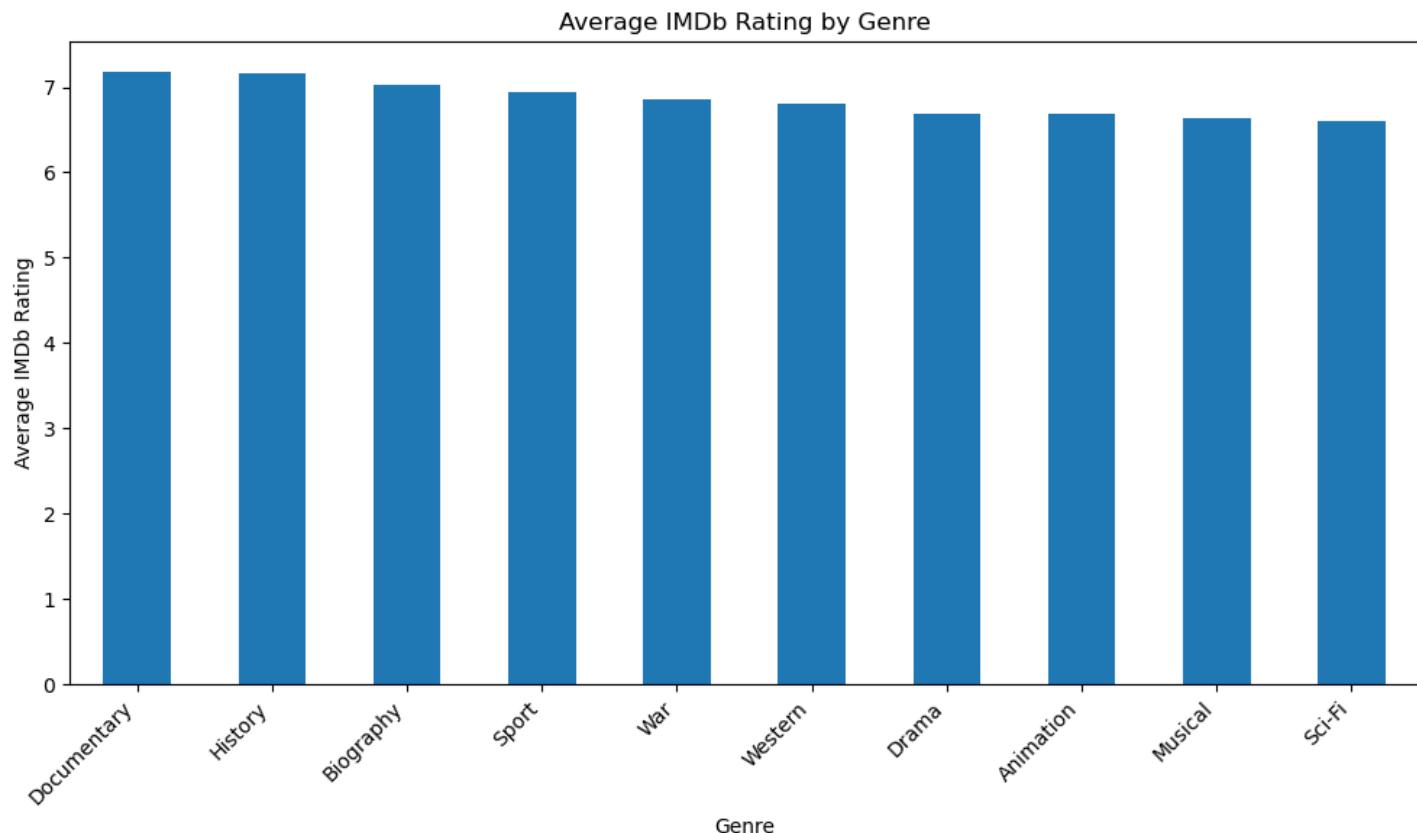
In [85]:

```
import matplotlib.pyplot as plt

top_genre_ratings = genre_ratings.head(10)

plt.figure(figsize=(10, 6))
top_genre_ratings.plot(kind='bar')

plt.title("Average IMDb Rating by Genre")
plt.xlabel("Genre")
plt.ylabel("Average IMDb Rating")
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



In [86]:

```
# show the distribution of ratings
```

```

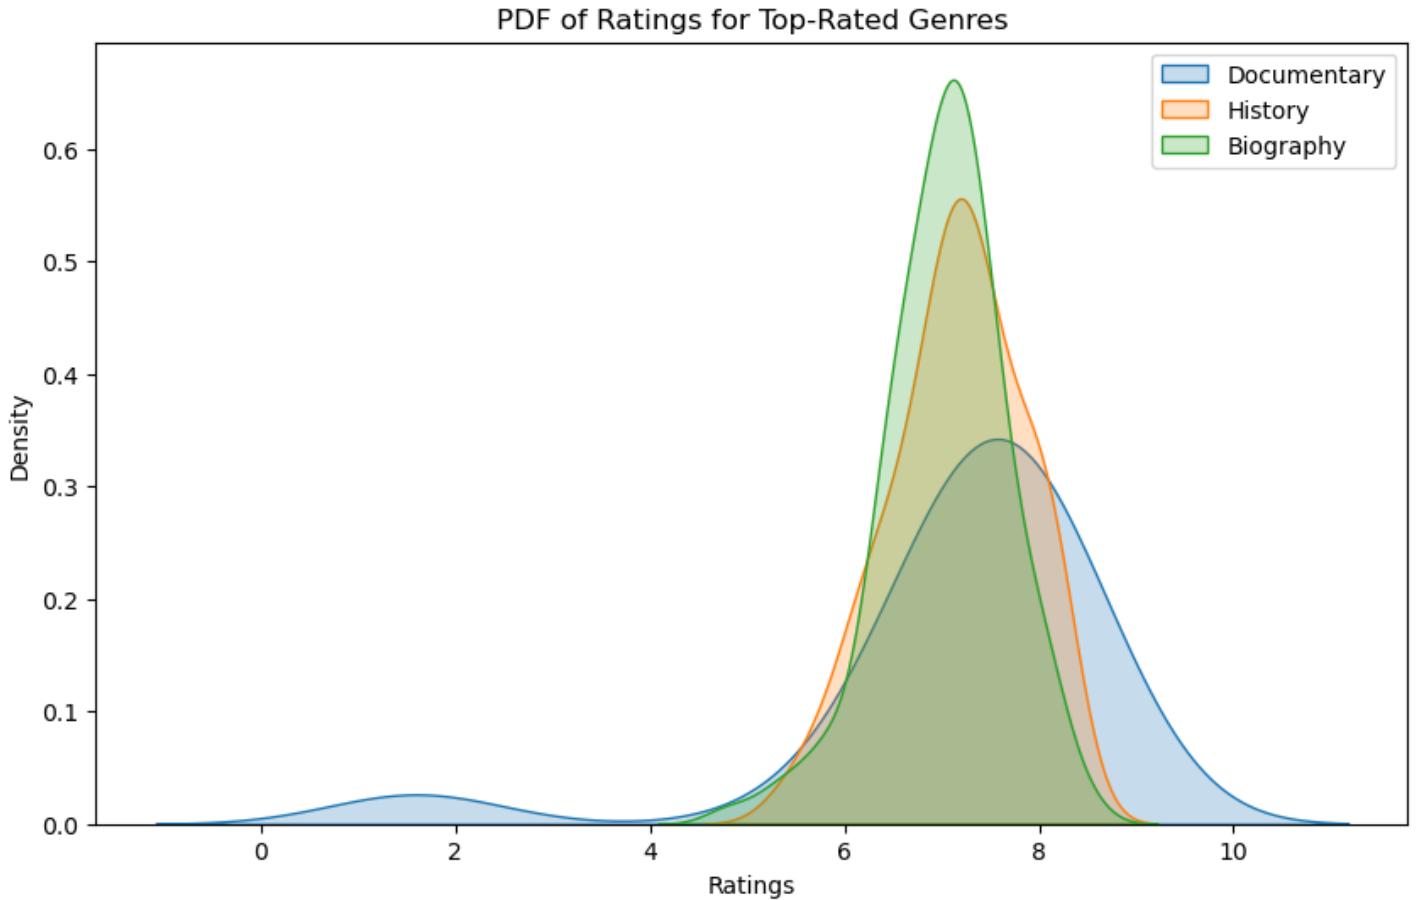
import seaborn as sns

top_genres = top_genre_ratings.index.tolist()

plt.figure(figsize=(10, 6))
for top_g in top_genres[:3]: # top 3 genres
    sns.kdeplot(
        themovie_df2[themovie_df2['genres'] == top_g]['averagerating'], label=top_g, fill=True
    )

plt.title("PDF of Ratings for Top-Rated Genres")
plt.xlabel("Ratings")
plt.ylabel("Density")
plt.legend()
plt.show()

```



In [87]:

```

from scipy.stats import f_oneway

samples = [
    themovie_df2[themovie_df2['genres'] == top_g]['averagerating']
    for top_g in top_genres[:5]
]

f_stat, p_value = f_oneway(*samples)

f_stat, p_value

```

Out [87]:

(0.6941202654410905, 0.5967399371566486)

Conclusion and Recommendation

Documentary, History and Biology movies receive the highest average ratings respectively.

A statistical test (ANOVA) on the genres,gave an F-statistic of 0.6941 and a p-value of 0.5967. Since the p-value

is above 0.05, there's no meaningful difference in average ratings between these genres. This Means good ratings can come from any genre.Hence, the studio should not primarily focus of ratings to measure the success of a movie.

Which genres generated the most profit Worldwide?

In [88]:

```
genre_revenue = (
    themovie_df2.groupby('genres')['profit'].mean().sort_values(ascending=False))

genre_revenue.head(20)
```

Out[88]:

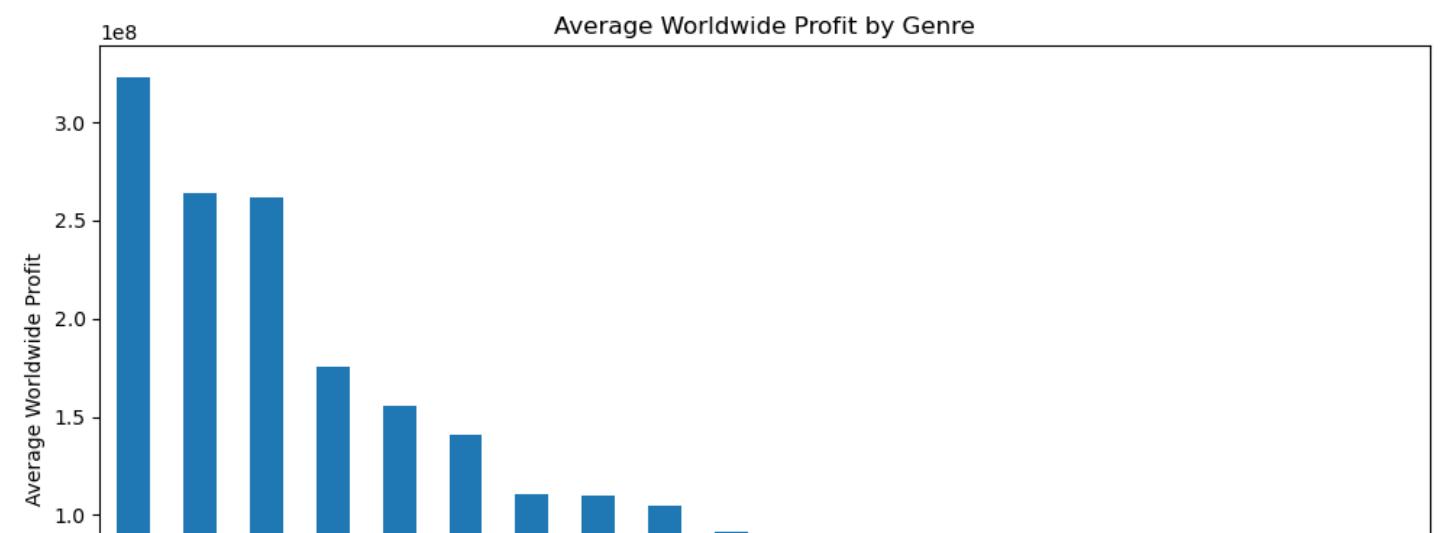
```
genres
Animation      3.227171e+08
Sci-Fi          2.641798e+08
Adventure       2.619090e+08
Action           1.753462e+08
Fantasy          1.555282e+08
Sport             1.404365e+08
Comedy            1.107014e+08
Musical           1.097501e+08
Family             1.046763e+08
Thriller          9.126921e+07
Mystery           7.510733e+07
Horror             6.993425e+07
Music              6.625840e+07
Western            6.078973e+07
Romance            5.978295e+07
Biography          5.497197e+07
Crime              5.352678e+07
Drama              5.143208e+07
History            5.054726e+07
War                 1.806302e+07
Name: profit, dtype: float64
```

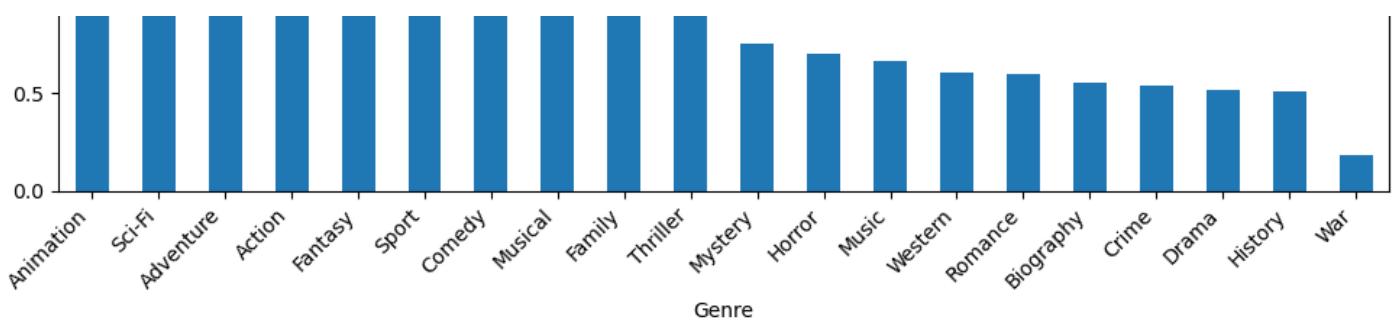
In [89]:

```
# Select top 10 genres for readability
top_genre_revenue = genre_revenue.head(20)

plt.figure(figsize=(10,6))
top_genre_revenue.plot(kind='bar')

plt.title("Average Worldwide Profit by Genre")
plt.xlabel("Genre")
plt.ylabel("Average Worldwide Profit")
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```





Animation, Sci-Fi and Adventure films generate the highest average worldwide profit, indicating strong global appeal in terms of revenue. The company should prioritize Action and Adventure films when launching its new movie studio.

Does Movie Rating Guarantees high Profits

In [90]:

```
correlation = themovie_df2['averagerating'].corr(themovie_df2['profit'])
correlation
```

Out[90]:

0.2567795120039979

Conclusion and Recomendation

This is a positive, but not strong. This means, they have a weak relationship. Higher-rated movies tend to make more profit, But ratings alone do not strongly determine profit. Therefore, good ratings help, but they do not guarantee high profit.

The success of a movie is influenced more by genre and market reach, and not ratings alone. Therefore, the studio should prioritize creating a movie based on genre and market reach but use rating as a secondary metrics.

Does release date (month) affect worldwide profit?

In [91]:

```
import scipy.stats as stats

# group by release month and calculate mean profit
monthly_profit = themovie_df2.groupby('release_month')['profit'].mean().reset_index()
monthly_profit = monthly_profit.sort_values('release_month')

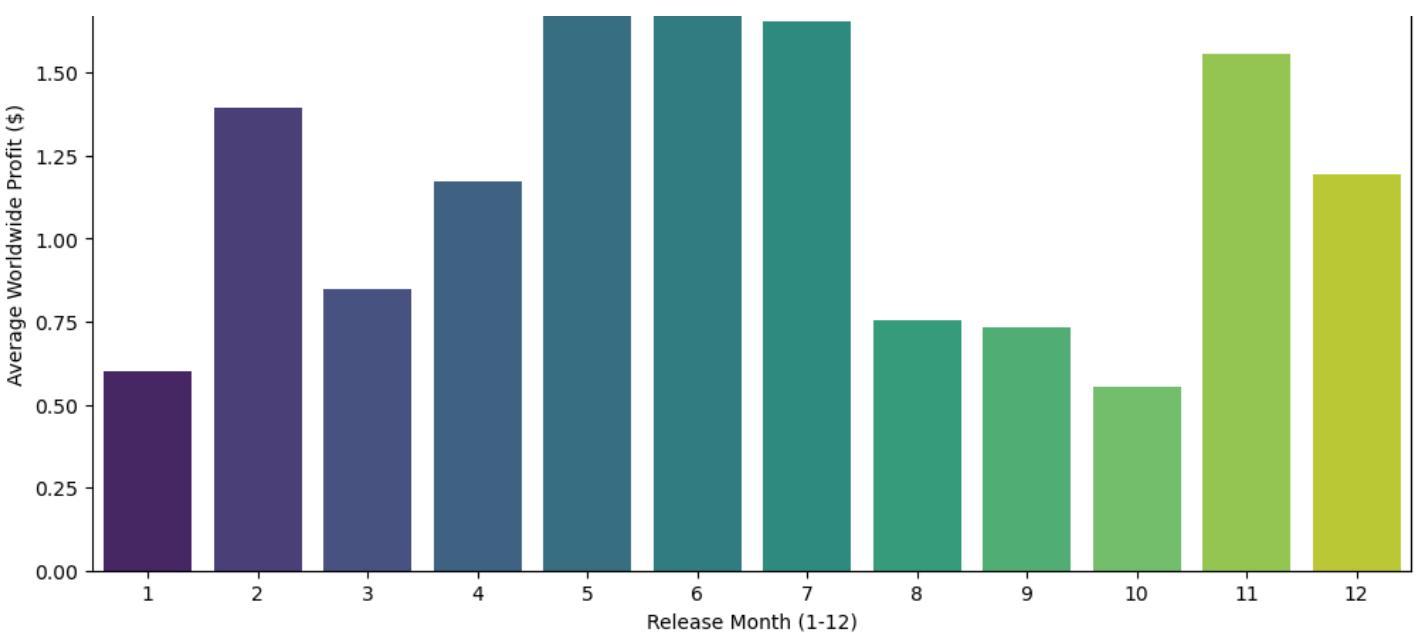
# Visualize
plt.figure(figsize=(12, 6))
sns.barplot(x='release_month', y='profit', data=monthly_profit, palette='viridis')
plt.title('Average Worldwide Profit by Release Month')
plt.xlabel('Release Month (1-12)')
plt.ylabel('Average Worldwide Profit ($)')
plt.show()
```

/tmp/ipykernel_21326/2919649226.py:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='release_month', y='profit', data=monthly_profit, palette='viridis')
```





Hypotheses

Null hypothesis: Worldwide profit is the same for all release months.

Alternative hypothesis: At least one release month has a different worldwide profit.

In [92]:

```
# Prepare data for ANOVA: list of profit arrays for each month
months = themovie_df2['release_month'].unique()
profit_by_month = [themovie_df2[themovie_df2['release_month'] == month]['profit'].values
for month in months]

# Perform ANOVA
f_stat, p_value = stats.f_oneway(*profit_by_month)

print(f"F-statistic: {f_stat}")
print(f"P-value: {p_value}")
```

F-statistic: 14.852492410836868
P-value: 1.4629188813546632e-28

From the observation we will reject the null hypothesis since our pvalue is less than 0.05 This means release month strongly affect movie profit. Time of release is important, So, Where releasing a movie the studio should consider the month which there are releasing the movie. They should release movies on month 2,5,6,7,11 and 12 since these months have the highest profit rate. They should avoid month 1,8,9 and 10.

What is the performance of Worldwide gross vs Domestic gross

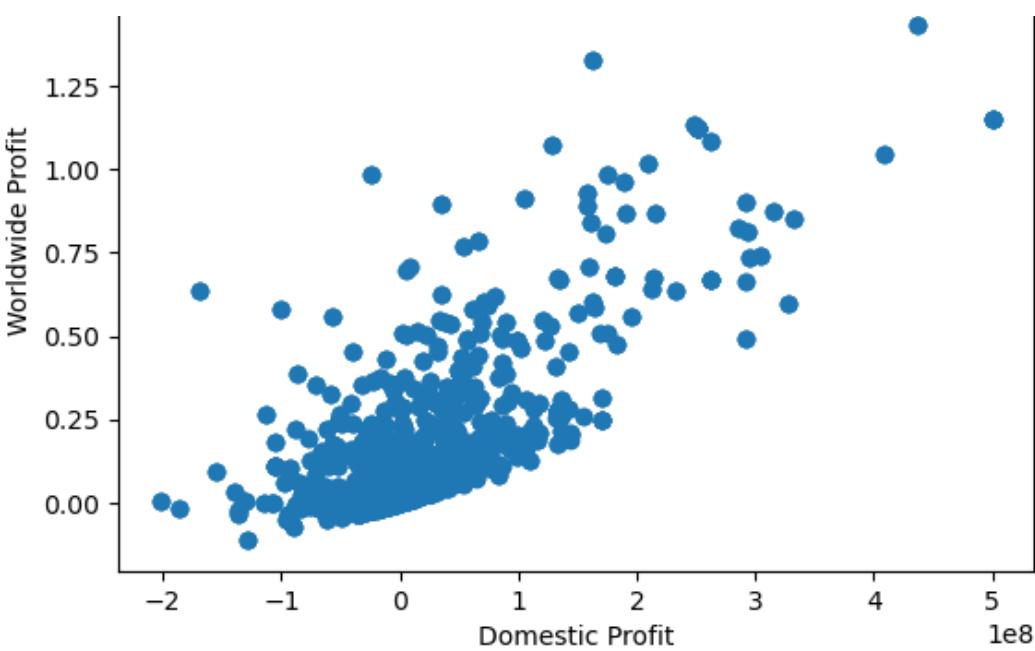
Goal: Determine whether movies earn more money internationally than domestically and how strongly the two are related.

In [93]:

```
plt.figure()
plt.scatter(themovie_df2["d_profit"], themovie_df2["profit"])
plt.xlabel("Domestic Profit")
plt.ylabel("Worldwide Profit")
plt.title("Relationship Between Domestic and Worldwide Profit")
plt.show()
```

1e9 Relationship Between Domestic and Worldwide Profit





In [94]:

```
corr, p_value = stats.pearsonr(themovie_df2["d_profit"], themovie_df2["profit"])

print("Correlation:", corr)
print("P-value:", p_value)
```

Correlation: 0.7274470666962394
P-value: 0.0

Conclusion and Recomendation

The correlation of 0.727447 indicates a very strong positive linear relationship between Domestic gross and worldwide gross p-value of 0 means the relationship is statistically significant so We reject the null hypothesis of no correlation This means if a movie makes more money in the domestic market, it is very likely to make more money worldwide as well.Hence Domestic performance of a movie should be used as a leading indicator for worldwide success of a movie.

Summary of Findings

Ratings vs Genre

Documentary, History, and Biology films receive the highest average ratings. However, an ANOVA test across genres shows no statistically significant difference in average ratings ($F = 0.6941$, $p = 0.5967$). This indicates that high ratings can occur in any genre, and ratings alone are not a reliable measure of a movie's success.

Ratings vs Profit

The relationship between movie ratings and profit is positive but weak. Higher-rated movies tend to earn more, but ratings alone do not strongly determine profitability. This suggests that while good ratings help, they do not guarantee high financial success.

Genre vs Worldwide Profit

Animation, Sci-Fi, and Adventure films generate the highest average worldwide profits, demonstrating strong global appeal. This indicates that genre choice has a greater impact on revenue than ratings. Action and Adventure films, in particular, should be prioritized by the studio.

Release Month vs Profit

Statistical analysis shows that release month has a significant effect on movie profit ($p < 0.05$). Movies released in February, May, June, July, November, and December perform best, while releases in January, August,

September, and October perform poorly. Timing is therefore a critical factor in maximizing profitability.

Domestic vs Worldwide Performance

There is a very strong positive correlation between domestic and worldwide gross ($r \approx 0.73$, $p \approx 0$). This means that movies that perform well domestically are very likely to perform well internationally. Domestic performance can be used as a leading indicator of global success.

Final Conclusion & Recommendation

Movie success is driven primarily by genre selection, release timing, and market reach, rather than ratings alone. While good ratings contribute positively, they should be treated as a secondary metric.

Based on Analysis we Recommend

1. **Prioritize Genre for Profitability:** Focus on producing Action, Adventure, Animation, and Sci-Fi movies, as these genres consistently generate the highest worldwide profits. Ratings are helpful but should not be the primary measure of expected success.
2. **Optimize Release Timing:** Schedule movie releases in months with historically high profits: February, May, June, July, November, and December. Avoid releasing in January, August, September, and October, which show lower profit potential.
3. **Use Domestic Performance as a Predictor:** Monitor domestic box office performance closely, as it is a strong indicator of worldwide success. Strong domestic earnings can help predict and maximize global revenue.
4. **Leverage Ratings Strategically:** While high ratings do not guarantee profit, they can enhance a movie's marketability. Use ratings as a secondary metric to guide marketing and audience targeting rather than as the primary decision factor for production or release strategy.